






# Intro to Stan

Jonah Gabry  
Columbia University






Stanislaw Ulam  
(1909–1984)



Stanislaw Ulam  
(1909–1984)

Monte Carlo  
Method



**Stanislaw Ulam  
(1909–1984)**

**H-Bomb**

**Monte Carlo  
Method**

What is Stan?

# What is Stan?

- Probabilistic **programming language** and **inference algorithms**



# What is Stan?

- Probabilistic **programming language** and **inference algorithms**
- Stan **program**
  - declares data and (constrained) parameter variables
  - defines log posterior (or penalized likelihood)

# What is Stan?

- Probabilistic **programming language** and **inference algorithms**
- Stan **program**
  - declares data and (constrained) parameter variables
  - defines log posterior (or penalized likelihood)
- Stan **inference**
  - MCMC for full Bayes
  - VB for approximate Bayes
  - Optimization for (penalized) MLE

Why Stan?

# Why Stan?

- Fit rich Bayesian statistical models

# Why Stan?

- Fit rich Bayesian statistical models
- Efficiency
  - HMC + NUTS
  - Compiled to C++

# Why Stan?

- Fit rich Bayesian statistical models
- Efficiency
  - HMC + NUTS
  - Compiled to C++
- Flexible domain specific language
  - Extensible
  - VB for approximate Bayes
  - Optimization for (penalized) MLE

# Why Stan?

- Fit rich Bayesian statistical models
- Efficiency
  - HMC + NUTS
  - Compiled to C++
- Flexible domain specific language
  - Extensible
  - VB for approximate Bayes
  - Optimization for (penalized) MLE
- Open source
  - BSD
  - CC-BY

Who is using Stan?



# Who is using Stan?

## Biological sciences

- clinical trials
- epidemiology
- genomics
- population ecology
- entomology
- ophthalmology
- neurology
- agriculture
- fisheries
- cancer biology

Who is using Stan?

Physical sciences

# Who is using Stan?

## Physical sciences

- astrophysics
  - LIGO gravitational wave observation
- molecular biology
- oceanography
- climatology

Who is using Stan?

Social sciences

# Who is using Stan?

## Social sciences

- population dynamics
- psycholinguistics
- social networks
- political science
- human development
- economics
  - textbook coming soon! (ish)

Who is using Stan?

More...

# Who is using Stan?

More...

- sports
- public health
- publishing
- finance
- pharma
- actuarial
- recommender systems
- educational testing
- materials engineering

Stan is many things



Stan is many things

**Math**

Stan is many things

**Math ← Language**

# Stan is many things

**Math** ← **Language** ← **Algorithms**

# Stan is many things

**Math ← Language ← Algorithms ← Services**


Stan is many things

**Math** ← **Stan**

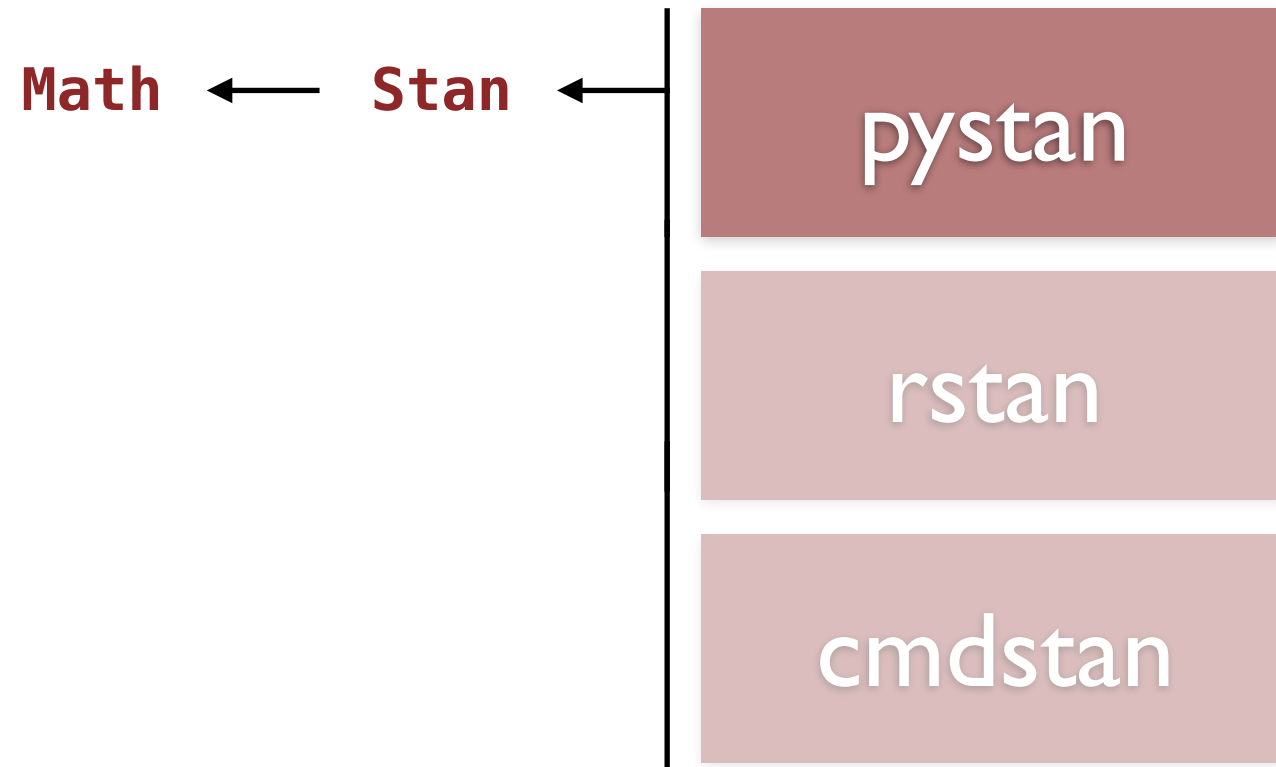
# Interfaces

**Math** ← **Stan**

# Interfaces




# Interfaces







# Interfaces




# Interfaces




# Interfaces




# Interfaces




# Interfaces



# Interfaces + Tools



# Interfaces + Tools







# Components of a Stan Program



Before we continue, install  
a few things:

# Before we continue, install a few things:

- If you haven't already, go to <https://github.com/jgabry/Bayes-Stan-Course>

# Before we continue, install a few things:

- If you haven't already, go to <https://github.com/jgabry/Bayes-Stan-Course>
- And then copy and run the code in the README file:

```
# install 'devtools' and then 'bayesplot'
if (!require("devtools")) install.packages("devtools")
library("devtools")
install_github("jgabry/bayesplot")

# also install 'shinystan'
install.packages("shinystan")
```

Posterior densities are specified in a comprehensive user-oriented probabilistic programming language.

$$p(q \mid \mathcal{D})$$

When writing a Stan program we always  
have three fundamental components

$$p(q \mid \mathcal{D})$$

When writing a Stan program we always  
have three fundamental components

$$p(q \mid \mathcal{D})$$

What are we conditioning on?

When writing a Stan program we always  
have three fundamental components

$$p(\boldsymbol{q} \mid \mathcal{D})$$

What are we conditioning on?

What are the parameters?



When writing a Stan program we always  
have three fundamental components

$$p(q \mid \mathcal{D})$$

What are we conditioning on?

What are the parameters?

How are they related?

We're now going to write a  
Stan program

We're now going to write a  
**Stan** program

- Open a new empty file in RStudio

# We're now going to write a Stan program

- Open a new empty file in RStudio
- Save it as **linear-regression.stan**



Stan programs are  
organized into *blocks*

Stan programs are  
organized into *blocks*

```
block name {
```

```
    block contents
```

```
}
```

Data



# Data

- Declare data types, sizes, and constraints

# Data

- Declare data types, sizes, and constraints
- Read from data source and constraints validated

# Data

- Declare data types, sizes, and constraints
- Read from data source and constraints validated
- Evaluated:
  - once

# Data

```
data {
```

```
// Dimensions
```

```
// Variables
```

```
}
```

# Data

```
data {  
  // Dimensions  
  int<lower=1> N;  
  
  // Variables  
  
}
```

# Data

```
data {
```

```
  // Dimensions
```

```
  int<lower=1> N;
```

```
  int<lower=1> K;
```

```
  // Variables
```

```
}
```

# Data

```
data {  
  // Dimensions  
  int<lower=1> N;  
  int<lower=1> K;  
  
  // Variables  
  matrix[N,K] X;  
  
}
```

# Data

```
data {  
  // Dimensions  
  int<lower=1> N;  
  int<lower=1> K;  
  
  // Variables  
  matrix[N,K] X;  
  vector[N] y;  
}
```



# Data

```
data {
```

```
  // Dimensions
```

```
  int<lower=1> N;
```

```
  int<lower=1> K;
```

```
  // Variables
```

```
  matrix[N,K] X;
```

```
  vector[N] y;
```

```
}
```

```
// single line comment
```

# Data

```
data {
```

```
  // Dimensions
```

```
  int<lower=1> N;
```

```
  int<lower=1> K;
```

```
  // Variables
```

```
  matrix[N,K] X;
```

```
  vector[N] y;
```

```
}
```

```
// single line comment
```

```
/* multiple lines of  
comments */
```

Parameters

# Parameters

- Declare parameter types, sizes, and constraints

# Parameters

- Declare parameter types, sizes, and constraints
- Transformations (under the hood) for constrained parameters

# Parameters

- Declare parameter types, sizes, and constraints
- Transformations (under the hood) for constrained parameters
- Evaluated:
  - every log prob evaluation

# Parameters

**parameters {**

**}**

# Parameters

```
parameters {  
    real alpha;  
  
}
```



# Parameters

**parameters {**

real alpha;

vector[K] beta;

**}**

# Parameters

```
parameters {  
    real alpha;  
    vector[K] beta;  
    real<lower=0> sigma;  
}
```

constraints *required* in  
**parameters** block

Model

# Model

- Statements defining the posterior density
  - log scale

# Model

- Statements defining the posterior density
  - log scale
- Evaluated:
  - every log prob evaluation

# Model

```
model {
```

```
}
```

# Model

```
model {  
  y ~ normal(X * beta + alpha, sigma);  
  
}
```

# Model

```
model {  
  y ~ normal(X * beta + alpha, sigma);  
  // priors (flat, uniform, if omitted)  
  
}
```



# Model

```
model {  
  y ~ normal(X * beta + alpha, sigma);  
  // priors (flat, uniform, if omitted)  
  
}
```

Why is the default automatically uniform?

# Model

```
model {  
  y ~ normal(X * beta + alpha, sigma);  
  // priors (flat, uniform, if omitted)  
  
}
```

Why is the default automatically uniform?

- $p(\theta) \propto 1$

# Model

```
model {  
  y ~ normal(X * beta + alpha, sigma);  
  // priors (flat, uniform, if omitted)  
  
}
```

Why is the default automatically uniform?

- $p(\theta) \propto 1$
- Nothing added to log prob

# Model

```
model {  
  y ~ normal(X * beta + alpha, sigma);  
  
  // priors (flat, uniform, if omitted)  
  alpha ~ normal(0, 10);  
  
}
```

Why is the default automatically uniform?

- $p(\theta) \propto 1$
- Nothing added to log prob

# Model

```
model {  
  y ~ normal(X * beta + alpha, sigma);  
  
  // priors (flat, uniform, if omitted)  
  alpha ~ normal(0, 10);  
  beta ~ normal(0, 10);  
  
}
```

Why is the default automatically uniform?

- $p(\theta) \propto 1$
- Nothing added to log prob

# Model

```
model {  
  y ~ normal(X * beta + alpha, sigma);  
  
  // priors (flat, uniform, if omitted)  
  alpha ~ normal(0, 10);  
  beta ~ normal(0, 10);  
  sigma ~ cauchy(0, 10);  
}
```

Why is the default automatically uniform?

- $p(\theta) \propto 1$
- Nothing added to log prob

# Generated Quantities

# Generated Quantities

- Declare and define derived variables
  - (P)RNGs, predictions, event probabilities, decision making



# Generated Quantities

- Declare and define derived variables
  - (P)RNGs, predictions, event probabilities, decision making
- Constraints validated

# Generated Quantities

- Declare and define derived variables
  - (P)RNGs, predictions, event probabilities, decision making
- Constraints validated
- Evaluated:
  - once per draw

# Generated Quantities

**generated quantities {**

**}**

# Generated Quantities

**generated quantities {**

vector[N] y\_rep;

**}**

# Generated Quantities

```
generated quantities {  
    vector[N] y_rep;  
    for (n in 1:N)  
  
}
```

# Generated Quantities

```
generated quantities {  
  vector[N] y_rep;  
  for (n in 1:N)  
    y_rep[n] =  
}
```

# Generated Quantities

```
generated quantities {  
  vector[N] y_rep;  
  for (n in 1:N)  
    y_rep[n] = normal_rng(  
}
```

# Generated Quantities

```
generated quantities {  
  vector[N] y_rep;  
  for (n in 1:N)  
    y_rep[n] = normal_rng(X[n] * beta + alpha,  
}
```



# Generated Quantities

```
generated quantities {
```

```
  vector[N] y_rep;
```

```
  for (n in 1:N)
```

```
    y_rep[n] = normal_rng(X[n] * beta + alpha, sigma);
```

```
}
```

**linear-regression.stan**

## linear-regression.stan

```
data {  
  int<lower=1> N;  
  int<lower=1> K;  
  matrix[N, K] X;  
  vector[N] y;  
}
```

## linear-regression.stan

```
data {  
  int<lower=1> N;  
  int<lower=1> K;  
  matrix[N, K] X;  
  vector[N] y;  
}  
parameters {  
  real alpha;  
  vector[K] beta;  
  real<lower=0> sigma;  
}
```

## linear-regression.stan

```
data {  
  int<lower=1> N;  
  int<lower=1> K;  
  matrix[N, K] X;  
  vector[N] y;  
}  
parameters {  
  real alpha;  
  vector[K] beta;  
  real<lower=0> sigma;  
}  
model {  
  y ~ normal(X * beta + alpha, sigma);  
}
```

## linear-regression.stan

```
data {  
  int<lower=1> N;  
  int<lower=1> K;  
  matrix[N, K] X;  
  vector[N] y;  
}  
parameters {  
  real alpha;  
  vector[K] beta;  
  real<lower=0> sigma;  
}  
model {  
  y ~ normal(X * beta + alpha, sigma);  
}  
generated quantities {  
  vector[N] y_rep;  
  for (n in 1:N)  
    y_rep[n] = normal_rng(X[n] * beta + alpha, sigma);  
}
```

# Transformed Data

# Transformed Data

- Declare and define transformed data variables



# Transformed Data

- Declare and define transformed data variables
- Constraints validated

# Transformed Data

- Declare and define transformed data variables
- Constraints validated
- Evaluated:
  - once (after **data**)

# Transformed Data

**transformed data {**

**}**

# Transformed Data

If we had declared  $X$  as  $K$  by  $N$  instead of  $N$  by  $K$   
we could transpose it here

**transformed data {**

**}**

# Transformed Data

If we had declared  $X$  as  $K$  by  $N$  instead of  $N$  by  $K$   
we could transpose it here

```
transformed data {  
    matrix[N,K] Xt;  
  
}
```

# Transformed Data

If we had declared  $X$  as  $K$  by  $N$  instead of  $N$  by  $K$  we could transpose it here

```
transformed data {  
    matrix[N,K] Xt;  
    Xt = X';  
}
```

# Transformed Data

If we had declared  $X$  as  $K$  by  $N$  instead of  $N$  by  $K$  we could transpose it here

```
transformed data {  
  matrix[N,K] Xt;  
  Xt = X'; // '=' for assignment since 2.10  
}
```

`Xt <- X'` deprecated  
in latest release  
(still allowed)

# Transformed Parameters



# Transformed Parameters

- Declare and define transformed parameter variables

# Transformed Parameters

- Declare and define transformed parameter variables
- Constraints validated

# Transformed Parameters

- Declare and define transformed parameter variables
- Constraints validated
- Evaluated:
  - every log prob evaluation

# Transformed Parameters

**transformed parameters {**

**}**

# Transformed Parameters

```
transformed parameters {  
    vector[N] eta;  
  
}
```

# Transformed Parameters

```
transformed parameters {  
    vector[N] eta;  
    eta = X * beta + alpha;  
}
```

# Functions

# Functions

- Declare and define functions to use in the body of the program



# Functions

- Declare and define functions to use in the body of the program
- Compiled with the model

# Functions

# Functions

**functions {**

```
vector linreg_DGP_rng(real alpha, vector beta,  
                      real sigma, matrix X) {  
  
    vector[rows(X)] y;  
    for (i in 1:rows(X))  
        y[i] = normal_rng(alpha + X[i] * beta, sigma);  
  
    return y;  
}
```

**}**