

# Section 3.

# Monte Carlo Methods

**Vincent Dorie**


New York University

**Part I**

# **Monte Carlo Integration**

# Monte Carlo Calculation of $\pi$

- Computing  $\pi = 3.14\dots$  via simulation is *the* textbook application of Monte Carlo methods.
- Generate points uniformly at random within the square
- Calculate proportion within circle ( $x^2 + y^2 < 1$ ) and multiply by square's area (4) to produce the area of the circle.
- This area is  $\pi$  (radius is 1, so area is  $\pi r^2 = \pi$ )



# Monte Carlo Calculation of $\pi$ (cont.)

- R code to calculate  $\pi$  with Monte Carlo simulation:

```
> x <- runif(1e6,-1,1)
```

```
> y <- runif(1e6,-1,1)
```

```
> prop_in_circle <- sum(x^2 + y^2 < 1) / 1e6
```

```
> 4 * prop_in_circle
```

```
[1] 3.144032
```

# Accuracy of Monte Carlo

- Monte Carlo is *not* an approximation!
- It can be made exact to within any  $\epsilon$
- Monte Carlo draws are i.i.d. by definition
- Central limit theorem: expected error decreases at rate of

$$\frac{1}{\sqrt{N}}$$


- 3 decimal places of accuracy with sample size 1e6
- Need 100× larger sample for each digit of accuracy

# General Monte Carlo Integration

- MC can calculate arbitrary definite integrals,

$$\int_a^b f(x) dx$$

- Let  $d$  upper bound  $f(x)$  in  $(a, b)$ ; tightness determines computational efficiency
- Then generate random points uniformly in the rectangle bounded by  $(a, b)$  and  $(0, d)$
- Multiply proportion of draws  $(x, y)$  where  $y < f(x)$  by area of rectangle,  $d \times (b - a)$ .
- Can be generalized to multiple dimensions in obvious way



# Expectations of Function of R.V.

- Suppose  $f(\theta)$  is a function of random variable vector  $\theta$
- Suppose the density of  $\theta$  is  $p(\theta)$
- Then  $f(\theta)$  is also random variable, with expectation

$$\mathbb{E}[f(\theta)] = \int f(\theta)p(\theta) \mathrm{d}\theta.$$

# Qol as Expectations

- Most Bayesian quantities of interest (Qol) are expectations over the posterior  $p(\theta | y)$  of functions  $f(\theta)$
- **Bayesian parameter estimation:**  $\hat{\theta}$ 
  - $f(\theta) = \theta$
  - $\hat{\theta} = \mathbb{E}[\theta | y]$  minimizes expected square error
- **Bayesian parameter (co)variance estimation:**  $\text{var}[\theta | y]$ 
  - $f(\theta) = (\theta - \mathbb{E}[\theta | y])^2$
- **Bayesian event probability:**  $\text{Pr}[A | y]$ 
  - $f(\theta) = \mathbb{I}[\theta \in A]$



# Expectations via Monte Carlo

- Generate draws  $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(M)}$  drawn from  $p(\theta)$
- Monte Carlo Estimator **plugs in average** for expectation:

$$\mathbb{E}[f(\theta) \mid y] \approx \frac{1}{M} \sum_{m=1}^M f(\theta^{(m)})$$

- Can be made **as accurate as desired**, because

$$\mathbb{E}[f(\theta)] = \lim_{M \rightarrow \infty} \frac{1}{M} \sum_{m=1}^M f(\theta^{(m)})$$

- *Reminder:* By CLT, error goes down as  $1 / \sqrt{M}$

**Part II**

# **Markov Chain Monte Carlo**

# Markov Chain Monte Carlo

- Standard Monte Carlo draws i.i.d. draws

$$\theta^{(1)}, \dots, \theta^{(M)}$$

according to a probability function  $p(\theta)$

- Drawing an i.i.d. sample is often impossible when dealing with complex densities like Bayesian posteriors  $p(\theta | y)$
- So we use Markov chain Monte Carlo (MCMC) in these cases and draw  $\theta^{(1)}, \dots, \theta^{(M)}$  from a Markov chain

# Markov Chains

- A Markov Chain is a sequence of random variables

$$\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(M)}$$

such that  $\theta^{(m)}$  only depends on  $\theta^{(m-1)}$ , i.e.,

$$p(\theta^{(m)} \mid y, \theta^{(1)}, \dots, \theta^{(m-1)}) = p(\theta^{(m)} \mid y, \theta^{(m-1)})$$

- Drawing  $\theta^{(1)}, \dots, \theta^{(M)}$  from a Markov chain according to  $p(\theta^{(m)} \mid \theta^{(m-1)}, y)$  is more tractable
- Require marginal of each draw,  $p(\theta^{(m)} \text{ mid } y)$ , to be equal to true posterior

# Applying MCMC

- Plug in just like ordinary (non-Markov chain) Monte Carlo
- Adjust standard errors for dependence in Markov chain

# MCMC for Posterior Mean

- Standard Bayesian estimator is posterior mean

$$\hat{\theta} = \int \theta p(\theta|y) d\theta$$

- Posterior mean minimizes expected square error
- Estimate is a conditional expectation

$$\hat{\theta} = \mathbb{E}[\theta | y]$$

- Compute by averaging

$$\hat{\theta} \approx \frac{1}{M} \sum_{m=1}^M \theta^{(m)}$$

# MCMC for Posterior Variance

- Posterior variance works the same way,

$$\mathbb{E}[(\theta - \mathbb{E}[\theta | y])^2 | y] \approx \frac{1}{M} \sum_{m=1}^M (\theta^{(m)} - \hat{\theta})^2$$

# MCMC for Event Probability

- Event probabilities are also expectations, e.g.,

$$\Pr[\theta_1 > \theta_2] = \mathbb{E}[\mathbb{I}[\theta_1 > \theta_2]] = \int_{\Theta} \mathbb{I}[\theta_1 > \theta_2] p(\theta | y) d\theta.$$

- Estimation via MCMC just another plug-in:

$$\Pr[\theta_1 > \theta_2] \approx \frac{1}{M} \sum_{m=1}^M \mathbb{I}[\theta_1^{(m)} > \theta_2^{(m)}]$$

- Again, can be made as accurate as necessary



# MCMC for Quantiles (incl. median)

- These are not expectations, but still plug in
- Alternative Bayesian estimator is posterior median
  - Posterior median minimizes expected absolute error
- Estimate as median draw of  $\theta^{(1)}, \dots, \theta^{(M)}$ 
  - just sort and take halfway value
  - e.g., Stan shows 50% point (or other quantiles)
- Other quantiles including interval bounds similar
  - estimate with quantile of draws
  - estimation error goes up in tail (based on fewer draws)

**Part III**

# **MCMC Algorithms**

# Random-Walk Metropolis

- We want to sample from a density  $p(\theta)$
- Take current state, add random perturbation
  - Proposal  $\theta^* = \theta^{(m)} + \epsilon$
- Accept new  $\theta$  with a probability that gives samples desired probability

Demo

- **Proposal** distribution determines behavior of chain

# Random-Walk Formal Def

- Draw random initial parameter vector  $\theta^{(1)}$  (in support)
- For  $m = 2, \dots, M$ 
  - Sample **proposal** from a (symmetric) jumping distribution, e.g.,

$$\theta^* \sim \text{MultiNormal}(\theta^{(m-1)}, \sigma I)$$

where  $I$  is the identity matrix

- Draw  $u \sim \text{Uniform}(0, 1)$  and set

$$\theta^{(m)} = \begin{cases} \theta^* & \text{if } u < \frac{p(\theta^* | y)}{p(\theta^{(m-1)} | y)} \\ \theta^{(m-1)} & \text{otherwise} \end{cases}$$

# Metropolis-Hastings

- Generalizes Metropolis to asymmetric proposals (i.e.  $J(\theta^* | \theta^{(m)}) \neq J(\theta^{(m)} | \theta^*)$ )
- Acceptance ratio is

$$\frac{J(\theta^{(m)} | \theta^*) \times p(\theta^* | y)}{J(\theta^* | \theta^{(m-1)}) \times p(\theta^{(m)} | y)}$$

where  $J$  is the proposal density

- i.e.,


$$\frac{\text{probability of being at } \theta^* \text{ and jumping to } \theta^{(m-1)}}{\text{probability of being at } \theta^{(m-1)} \text{ and jumping to } \theta^*}$$

## Metropolis-Hastings (cont.)


- General form ensures equilibrium by maintaining *detailed balance*
- Like Metropolis, only requires ratios
- Many algorithms involve a Metropolis-Hastings “correction”
  - Including vanilla HMC and RHMC and ensemble samplers
- **Detailed Balance** & **Reversibility** formal conditions to guarantee *convergence* of chain to desired distribution

# Optimal Proposal Scale?


- Proposal scale  $\sigma$  is a free; too low or high is inefficient



(a) Proposal variance too large



(b) Proposal variance too small



(c) Proposal variance approximately optimised

- Traceplots* show parameter value on  $y$  axis, iterations on  $x$
- Empirical tuning problem; theoretical optima exist for some cases


# Convergence

- Markov chains depend on previous state
- Construct so that stationary (long-run) state is target
- What about starting state?
- Imagine a signal problem on the MTA
  - A trains on the F line to Brooklyn, C service suspended...
  - Substantial delays at W 4th St
  - Eventually... service returns to normal
- May take many iterations for Markov chain to reach equilibrium



# Convergence: Example

- Test for convergence: start multiple Markov chains at diffuse points



- Four chains with different starting points
  - *Left*: 50 iterations
  - *Center*: 1000 iterations
  - *Right*: Draws from second half of each chain

Gelman et al., *Bayesian Data Analysis*

# Convergence Diagnostic ( $\hat{R}$ )

- Gelman & Rubin recommend  $M$  chains of  $N$  draws with **diffuse initializations**
- Measure that each chain has same posterior mean and variance
- If not, may be stuck in multiple modes or just not converged yet
- Define statistic  $\hat{R}$  of chains s.t. **at convergence**,  $\hat{R} \rightarrow 1$ 
  - $\hat{R} \gg 1$  implies non-convergence
  - $\hat{R} \approx 1$  **does not guarantee convergence**
  - Only measures marginals


# Effective Sample Size

- Markov chains typically display autocorrelation in the series of draws  $\theta^{(1)}, \dots, \theta^{(m)}$
- Without i.i.d. draws, central limit theorem *does not apply*
- Effective sample size  $N_{\text{eff}}$  divides out autocorrelation
- $N_{\text{eff}}$  must be estimated from sample
- Estimation accuracy proportional to


$$\frac{1}{\sqrt{N_{\text{eff}}}}$$

# Reducing Posterior Correlation


- Tuning algorithm parameters to ensure good mixing
- Recall Metropolis traceplots of Roberts and Rosenthal:



(a) Proposal variance too large



(b) Proposal variance too small



(c) Proposal variance approximately optimised

- Good jump scale  $\sigma$  produces good mixing and high  $N_{\text{eff}}$

# Gibbs Sampling

- Draw random initial parameter vector  $\theta^{(1)}$  (in support)
- For  $m = 2, \dots, M$ :
  - For  $n = 1, \dots, N$ :
    - \* draw  $\theta_n^{(m)}$  according to conditional

$$p(\theta_n \mid \theta_1^{(m)}, \dots, \theta_{n-1}^{(m)}, \theta_{n+1}^{(m-1)}, \dots, \theta_N^{(m-1)}, y).$$

- e.g, with  $\theta = (\theta_1, \theta_2, \theta_3)$ :
  - draw  $\theta_1^{(m)}$  according to  $p(\theta_1 \mid \theta_2^{(m-1)}, \theta_3^{(m-1)}, y)$
  - draw  $\theta_2^{(m)}$  according to  $p(\theta_2 \mid \theta_1^{(m)}, \theta_3^{(m-1)}, y)$
  - draw  $\theta_3^{(m)}$  according to  $p(\theta_3 \mid \theta_1^{(m)}, \theta_2^{(m)}, y)$

# Generalized Gibbs

- “Proper” Gibbs requires conditional Monte Carlo draws
  - typically works only for conjugate priors
- In general case, may need to use less efficient conditional draws
  - Slice sampling is a popular general technique that works for discrete or continuous  $\theta_n$  (JAGS)
  - Adaptive rejection sampling is another alternative (BUGS)
  - Very difficult in more than one or two dimensions

# Sampling Efficiency

- We care only about  $N_{\text{eff}}$  per second
- Decompose into
  1. Iterations per second
  2. Effective sample size per iteration
- Gibbs and Metropolis have high iterations per second (especially Metropolis)
- But they have low effective sample size per iteration (especially Metropolis)
- Both are particular weak when there is high correlation among the parameters in the posterior

# Hamiltonian Monte Carlo & NUTS

- Slower iterations per second than Gibbs or Metropolis
  - Much higher effective sample size per iteration for complex posteriors (i.e., high curvature and correlation)
  - Overall, much higher  $N_{\text{eff}}$  per second
- 
- Details in the next talk ...
  - Along with details of how Stan implements HMC and NUTS



# Why Stan is Great

- Writing a sampler for model:

$$y \mid \beta, \sigma^2 \sim \text{Normal}(x\beta, \sigma_y^2),$$

$$\beta \sim \text{Normal}(0, \sigma_\beta^2),$$

$$\sigma \sim \text{Half-Cauchy}.$$

**The End (Section 3)**

**Part I**

# **What Stan Does**

# Full Bayes: No-U-Turn Sampler


- Adaptive **Hamiltonian Monte Carlo** (HMC)
  - **Potential Energy**: negative log posterior
  - **Kinetic Energy**: random standard normal per iteration
- Adaptation **during warmup**
  - step size adapted to target total acceptance rate
  - mass matrix estimated with regularization
- Adaptation **during sampling**
  - simulate forward and backward in time until U-turn
- **Slice sample** along path

(Hoffman and Gelman 2011, 2014)

# Sample: Hamiltonian Flow

- Generate random **kinetic energy**
  - random  $\text{Normal}(0, 1)$  in each parameter
- Use negative log posterior as **potential energy**
- Hamiltonian is kinetic plus potential energy
- **Leapfrog Integration**: for *fixed* stepsize (time discretization), number of steps (total time), and mass matrix,
  - update momentum half-step based on potential (gradient)
  - update position full step based on momentum
  - update momentum half-step based on potential
- Numerical solution of Hamilton's first-order version of Newton's second-order diff-eqs of motion (force = mass  $\times$  acceleration)

# Sample: Leapfrog Example




- Trajectory of 25 leapfrog steps for correlated 2D normal (ellipses at 1 sd from mean), stepsize of 0.25, initial state of  $(-1, 1)$ , and initial momentum of  $(-1.5, -1.55)$ .

Radford Neal (2013) MCMC using Hamiltonian Dynamics. In *Handbook of MCMC*. (free online at <http://www.mcmchandbook.net/index.html>)

# Sample: No-U-Turn Sampler (NUTS)

- Adapts Hamiltonian simulation time
  - goal to maximize mixing, maintaining detailed balance
  - too short devolves to random walk
  - too long does extra work (i.e., orbits)
- For exponentially increasing number of steps up to max
  - Randomly choose to extend forward or backward in time
  - Move forward or backward in time number of steps
    - \* stop if any subtree (size 2, 4, 8, ...) makes U-turn
    - \* remove all current steps if subtree U-turns (not ends)
- Randomly select param with density above slice (or reject)

# Sample: NUTS Binary Tree




- Example of repeated doubling building binary tree forward and backward in time until U-turn.

Hoffman and Gelman. 2014. The No-U-Turn Sampler. *JMLR*. (free online at <http://jmlr.org/papers/v15/hoffman14a.html>)



# Sample: NUTS U-Turn



- Example of trajectory from one iteration of NUTS.
- Blue ellipse is contour of 2D normal.
- Black circles are leapfrog steps.
- Solid red circles excluded below slice
- U-turn made with blue and magenta arrows
- Red crossed circles excluded for detailed balance

# Sample: HMC/NUTS Warmup

- Estimate stepsize
  - too small requires too many leapfrog steps
  - too large induces numerical inaccuracy
  - need to balance
- Estimate mass matrix
  - Diagonal accounts for parameter scales
  - Dense optionally accounts for rotation


# Sample: Warmup (cont.)

- Initialize unconstrained parameters as for optimization
- For exponentially increasing block sizes
  - for each iteration in block
    - \* generate random kinetic energy
    - \* simulate Hamiltonian flow (HMC fixed time, NUTS adapts)
    - \* choose next state (Metropolis for HMC, slice for NUTS)
  - update regularized point estimate of mass matrix
    - \* use parameter draws from current block
    - \* shrink diagonal toward unit; dense toward diagonal
  - tune stepsize (line search) for target acceptance rate

# Sample: HMC/NUTS Sampling


- Fix stepsize and mass matrix
- For sampling iterations
  - generate random kinetic energy
  - simulate Hamiltonian flow
  - apply Metropolis accept/reject (HMC) or slice (NUTS)

# NUTS vs. Gibbs and Metropolis



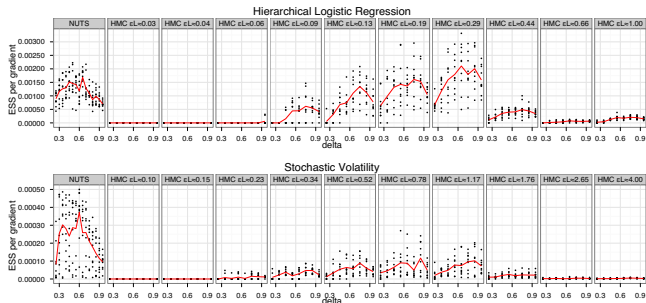
- Two dimensions of highly correlated 250-dim normal
- **1,000,000 draws** from Metropolis and Gibbs (thin to 1000)
- **1000 draws** from NUTS; 1000 independent draws

# NUTS vs. Basic HMC



- 250-D normal and logistic regression models
- Vertical axis is effective sample size per sample (bigger better)
- Left) NUTS; Right) HMC with increasing  $t = \epsilon L$

# NUTS vs. Basic HMC II



- Hierarchical logistic regression and stochastic volatility
- Simulation time  $t$  is  $\epsilon L$ , step size ( $\epsilon$ ) times number of steps ( $L$ )
- NUTS can beat optimally tuned HMC (latter very expensive)

**Part III**

# **Under Stan's Hood**



# Euclidean Hamiltonian

- **Phase space:**  $q$  position (parameters);  $p$  momentum
- **Posterior density:**  $\pi(q)$
- **Mass matrix:**  $M$
- **Potential energy:**  $V(q) = -\log \pi(q)$
- **Kinetic energy:**  $T(p) = \frac{1}{2} p^\top M^{-1} p$
- **Hamiltonian:**  $H(p, q) = V(q) + T(p)$
- **Diff eqs:**

$$\frac{dq}{dt} = + \frac{\partial H}{\partial p} \qquad \frac{dp}{dt} = - \frac{\partial H}{\partial q}$$

# Leapfrog Integrator Steps

- Solves Hamilton's equations by **simulating dynamics** (symplectic [volume preserving];  $\epsilon^3$  error per step,  $\epsilon^2$  total error)
- Given: **step size**  $\epsilon$ , **mass matrix**  $M$ , **parameters**  $q$
- **Initialize kinetic** energy,  $p \sim \text{Normal}(0, \mathbf{I})$
- **Repeat** for  $L$  leapfrog steps:

$$p \leftarrow p - \frac{\epsilon}{2} \frac{\partial V(q)}{\partial q} \quad [\text{half step in momentum}]$$

$$q \leftarrow q + \epsilon M^{-1} p \quad [\text{full step in position}]$$



$$p \leftarrow p - \frac{\epsilon}{2} \frac{\partial V(q)}{\partial q} \quad [\text{half step in momentum}]$$

**Part IV**

# **Mixutre Models**

# Mixture Models

- Class of models with observations coming from than one distribution, membership unknown




# Mixture Model Density

- Depends on a **latent** class membership

$$y_i \mid z_i \sim \begin{cases} p_1(y_i) & \text{if } z_i = 1 \\ p_2(y_i) & \text{if } z_i = 2 \\ \vdots & \\ p_K(y_i) & \text{if } z_i = K \end{cases},$$
$$z_i \sim \text{Categorical}(\lambda).$$

$$\lambda \geq 0, \sum_{k=1}^K \lambda_k = 1$$

# Example Density



- $y_i | z_i \sim \text{Normal}(4.5, 0.1)$  if  $z_i = 1$ ,
- $y_i | z_i \sim \text{Normal}(2, 2)$  if  $z_i = 2$ ,
- $y_i | z_i \sim \text{Normal}(7, 1.25)$  if  $z_i = 3$ ,
- $\lambda = (0.2, 0.5, 0.3)$

```
data {  
  int<lower = 0> N;  
  int<lower = 0> K;  
  real y[N];  
}  
parameters {  
  simplex[K] lambda;  
  int z[K];  
  
  real[K] mu;  
  real<lower = 0> sigma[K];  
}  
model {  
  z ~ categorical(lambda);  
  for (n in 1:n)  
    y[n] ~ normal(mu[z[n]], sigma[z[n]])  
}
```

```

data {
  int<lower = 0> N;
  int<lower = 0> K;
  real y[N];
}
parameters {
  simplex[K] lambda;
  int z[K];

  real[K] mu;
  real<lower = 0> sigma[K];
}
model {
  z ~ categorical(lambda);
  for (n in 1:n)
    y[n] ~ normal(mu[z[n]], sigma[z[n]])
}

```

integer parameters or transformed parameters are not allowed;  
 found declared type int, parameter name=z Problem with  
 declaration.



## Integrate/Sum Out $z$

$$p(y_i, z_i) = [I[z_i = 1]p_1(y_i) + I[z_i = 2]p_2(y_i) + \cdots + I[z_i = K]p_K(y_i)] \prod_{k=1}^K \lambda_k^{I[z_i=k]}.$$

$$\begin{aligned} \sum_{z_i=1}^K p(y_i, z_i) &= \lambda_1 p_1(y_i) + \lambda_2 p_2(y_i) + \cdots + \lambda_K p_K(y_i), \\ &= p(y_i). \end{aligned}$$

Multiply together to get likelihood:

$$p(y) = \prod_{i=1}^N [\lambda_1 p_1(y_i) + \lambda_2 p_2(y_i) + \cdots + \lambda_K p_K(y_i)].$$

# Log Problem

- Operate on log probability as multiplication leads to underflow
  - small #  $\times$  small #  $\times$  ...
  - $\log(\text{small \#}) + \log(\text{small \#}) + \dots$
- Log-likelihood is unwieldy:

$$\log p(y) = \sum_{i=1}^N \log [\lambda_1 p_1(y_i) + \lambda_2 p_2(y_i) + \dots + \lambda_K p_K(y_i)].$$

log\_sum\_exp